
Warranty GENE Warehouse Wireframe

A grain-declared structured database over the CDK Drive
and PACCAR warranty extract set — v0.1

PROGRAMME

Hawkins Twin Platform · CDK Twin · Peterbilt Atlantic

DOCUMENT

EgD-HAW-CDK-WH-001 r3

EXTRACT SET

2026-09-17 — nine loaded source files, six logged out of scope, 2,991 units, 7,794 repair orders

JOIN SPINE

VIN8 — manufacturer registration to dealer repair order

REPOSITORY

VEglyphDesign/eve-hawkins-cdk-twin · warehouse/

PUBLIC SURFACE

eveglyphdesign.github.io/eve-hawkins-cdk-twin/warehouse/

CUSTODY

Schema and rules published · dealer payload never committed

EVEGLYPHDESIGN · CONTROLLED COPY

This directory is the structured database wireframe for the Peterbilt Atlantic warranty lane. It converts a pile of ad-hoc CDK Drive report exports and PACCAR warranty registration reports into one grain-declared star schema with a single join spine, so that every downstream worklist, dashboard and alert reads one shape and no other.

It is a **wireframe**: schema, grain, keys, transforms, marts and data-quality rules.

Scope — the truck business, and only the truck business

This model covers **Peterbilt Atlantic's truck lane**: the eight CDK service accounts under accounting account `PNB-A` , the PACCAR warranty registrations behind those units, and the customer and advisor masters that serve them. Every unit in it is a Class 8 or medium-duty Peterbilt carrying PACCAR coverage — MX-11 and MX-13 engine, aftertreatment, PACCAR TX-12 transmission and clutch, harness, tow, and base vehicle.

`TRPDT-S` belongs here. It is **TRP Dartmouth** — PACCAR's all-makes truck parts and service brand — and its customers are fuel haulers and transport fleets, not powersports.


The **BRP / powersports lane is deliberately out of scope**: the Advantage Plus and Canada General service contracts, the Warranty on Demand claim records, the BRP warranty policy guide, and the Torque Motorsports retention reports. They are present in the extract directory, and the loader logs each one as a `SCOPE_EXCLUDED` row in `dq_exception` so the exclusion is auditable rather than silent. That lane is a different VIN space with a different policy owner and belongs in its own model. Joining it onto a truck VIN8 spine is how one machine's coverage gets attached to another's.

Custody

No dealer rows are committed to this repository. This repo holds the index, the rules and the schema — never the payload. The loader reads a dealer-owned extract directory and writes the database file outside the repository tree. The only published artefact derived from the data is `profile/2026-09-17-extract-profile.json` , which contains row counts and aggregates and no customer, VIN or contract row.

The join spine — VIN8

The single most important finding in this extract set: **VIN8 joins the manufacturer record to the dealer record**. CDK's `Vehicle ID` on the repair order carries the 8-character VIN suffix, and every PACCAR warranty-registration row is keyed on `VIN8` .



Measured yield on the 2026-09-17 set:

	UNITS
Distinct units in <code>dim_unit</code>	2,991
Warranty-registered (PACCAR side)	1,060
Service-known (CDK RO side)	2,267
Both sides — fully reconciled	336

The secondary key is the customer: CDK's legacy Pick `NAME-FILE` encodes the customer number as `<store_prefix>*<customer_number>` , which resolves 818 of the 821 customers appearing on repair orders back to the

customer master. Registered-owner name is *not* a usable join — only 65 of 534 manufacturer-side owner names match a CDK RO customer name exactly, which is exactly why the model joins on VIN8 and treats names as attributes.

Model

Six dimensions, two facts, two reconciliation aggregates, one quarantine table, eleven marts.

Dimensions — `dim_unit` (VIN8), `dim_customer`, `dim_rooftop`, `dim_warranty_option`, `dim_service_advisor`, `dim_date`.

Facts — `fact_warranty_registration` (grain: VIN8 × warranty option, manufacturer truth), `fact_repair_order` (grain: one RO, open and closed unioned under `ro_status`, dealer truth).

Aggregates — `agg_customer_service_period`, `agg_customer_total_sales`. These are tie-out instruments only. They are never a scoring input; a pre-aggregated total cannot be attributed to a VIN and must not be allowed to look as if it can.

Quarantine — `dq_exception`. Nothing is silently dropped. Every unjoinable vehicle ID, unparseable date and unknown customer lands as a row with a rule code and a severity.

Full ERD in `ERD.md`. DDL in `ddl/wireframe.sql`. Field-level lineage for all 40 mapped columns in `mapping/extract_to_model.csv`.

Marts — the only shapes GENE may read

VIEW	GRAIN	PURPOSE
<code>v_unit_coverage_state</code>	VIN8 × option	coverage window state against a run date, both axes
<code>v_unit_service_profile</code>	VIN8	RO count, last visit, max odometer, lifetime sales
<code>v_rich_target</code>	VIN8	active coverage plus a reachability class
<code>v_whitespace_registered_never_served</code>	VIN8	registered in territory, never in the bays
<code>v_coverage_expiring_180d</code>	VIN8	expiry pressure inside two quarters
<code>v_served_unregistered</code>	VIN8	served here, no manufacturer registration on file
<code>v_rooftop_extract_coverage</code>	service account	which of the eight rooftops actually have history
<code>v_unit_mileage_rate</code>	VIN8	measured miles/day from the unit's own odometer readings
<code>v_system_weight</code>	system code	diagnostic exposure weight per covered system (policy)
<code>v_pre_ro_priority</code>	VIN8	perishability on the binding axis, exposure, cost of contact
<code>v_diagnosis_queue</code>	VIN8	the banded pre-RO work queue, P1 to P6

`v_rich_target` classes every unit with live coverage by how the dealership can actually reach it:

REACH CLASS	UNITS	EXPIRING WITHIN 180 DAYS
<code>ON_SITE_NOW</code> — on an open RO today	60	14
<code>KNOWN_CUSTOMER</code> — has RO history here	205	35
<code>REGISTERED_NEVER_SERVICED</code> — registered in territory, never serviced	508	68

What the wireframe already shows

- **773 units carry live manufacturer coverage**, holding 1,802 active coverages — 1,236 engine, 273 aftertreatment, 257 vehicle, plus emissions, electrical, towing and cab-clothing lines.
- **60 units with active coverage are on an open repair order right now.** One is 1 day from time expiry. This is the pre-RO decision window GENE exists to serve.
- **508 registered units have no service history in the extract set — an upper bound, not a count.** Closed history covers three of the group's **eight** service accounts, so a unit serviced only at PNBMS , PNBFS , PQSP-S , PQC-S or TRPDT-S appears here as whitespace and is not. Treat this view as a re-extract requirement, not a campaign list. See Extract coverage .
- **170 of 650 time-active coverages are already past their mileage ceiling.** Coverage is a two-axis test — months and miles — and a time-only filter would mis-flag every one of them. v_unit_coverage_state therefore exposes is_time_active and miles_remaining_est as separate columns and refuses to collapse them into one boolean.
- Warranty work is **14.7% of service sales** in the extract window, against \$30.0M customer-pay. The wireframe's job is to move that share deliberately rather than by accident.
- Dealer-side history in scope: 7,534 closed repair orders, \$12.11M total sales, \$5.90M gross profit, 2024-10-01 through 2026-09-16 — but drawn from only **three of eight** service accounts, all under one accounting account (PNB-A). Those figures are a floor for the group, not the group.

Extract coverage — eight rooftops, three histories

dim_rooftop is seeded from the union of the open and closed RO extracts, because the two extracts do not cover the same dealership. Eight service accounts appear; only three carry closed history. Labels are left UNVERIFIED until the dealership confirms them — the codes are not decoded by guesswork.

SERVICE ACCOUNT	CLOSED ROS	OPEN ROS	OPEN WIP SALES	CLOSED HISTORY PRESENT
PBNS-S	3,036	33	\$48,089.69	yes
PBDT-S	2,647	31	\$64,914.32	yes
PNBDL-S	1,863	27	\$63,236.60	yes
PNBM-S	0	60	\$277,334.33	no
PNBF-S	0	41	\$153,227.33	no
TRPDT-S (TRP Dartmouth)	0	30	\$68,495.06	no
PQSP-S	0	20	\$133,463.93	no
PQC-S	0	18	\$23,549.02	no

The five uncovered accounts hold **169 of 260 open repair orders — 65% of current work in progress, and \$656,069.67 of \$832,310.28 in open sales.** The rooftops missing from history are the busiest ones in the present. 125 distinct units are on an open RO at an uncovered rooftop; 46 of them are warranty-registered and 31 already qualify for v_rich_target .

Every uncovered account is recorded as a R00FTOP_CLOSED_HISTORY_MISSING row in dq_exception and is queryable through v_rooftop_extract_coverage . Ask 1 of the next extract request is closed ROs for all eight service accounts over the same window.

Opportunistic-diagnosis priority — the pre-RO queue

`v_diagnosis_queue` ranks the 703 units whose coverage is alive on **both** axes. 70 units with live months but a spent mileage ceiling are excluded on purpose. Bands are instructions to a service writer; the score only orders work inside a band.

BAND	UNITS	MEANING
P1 IN THE BAY, CLOSING	5	on an open RO now, 30 days or less on the binding axis
P2 IN THE BAY	42	on an open RO now, coverage live
P3 CALL NOW, CLOSING	28	RO history here, 90 days or less
P4 CALL, SCHEDULE	120	RO history here, coverage live
P5 COLD, CLOSING	38	registered only, 90 days or less
P6 COLD, CAMPAIGN	470	registered only, coverage live

Score — 0-100. Perishability on the binding axis 60, diagnostic exposure by covered system 25, cost of contact 15. System weights (`v_system_weight`) are ENG 4, EMC 4, A/T 3, VEH 2, ELEC 2, TOW 1, CLTH 1. They are policy, held in a view so they can be argued with rather than buried in a formula.

The mileage clock. `v_unit_mileage_rate` derives miles per day from each unit's own odometer readings — two or more readings, 60 days or more apart, positive gain. Rates above 1,200 miles/day are flagged `SUSPECT_ODOMETER` and never used to date an expiry. **17 units read as safe on months and sit inside 90 days on miles; eleven of those show more than six months left on the calendar, one shows 2,741 days.** A months-only expiry report ranks the most urgent trucks in the file as low priority.

Where no rate can be measured the queue falls back to the months axis and says so via `rate_status = 'NO_RATE'`. That is the whole cold lane — and 12 of the 13 covered units on an open RO at `PNBM-S` right now, because Moncton's closed history was never extracted. Their priority is understated for an extract reason, not a business one.

Contact reach — 188 of the 195 warm units (P1–P4) have a phone number on file, 167 have an email. The 508 cold units have no phone at all: registered owner name and city only, resolving to 313 distinct owners across 241 cities, which makes that lane a territory-rep list rather than a call list. Eleven owners hold four or more covered units each.

No claim-value estimate per row. Without RO line detail there are no labour operation codes or SRT hours, so the queue says which trucks are worth diagnosing and cannot say what the diagnosis is likely to be worth.

Outreach workbook

`bin/build_outreach_workbook.py` renders bands P1–P4 to an Excel workbook, one sheet per band, every ranking signal as a column alongside the contact details held for the customer on the unit's most recent RO, plus a coverage-detail sheet, a **CDK vs PACCAR** sheet and a method sheet.

The `CDK vs PACCAR` sheet is the two-sided reconciliation: one row for every one of the 2,991 units in the model, the manufacturer extract's account of the unit set beside the dealer system's account of it, and for any unit outside P1–P4 the reason it is outside. `match_status` splits into **BOTH SIDES 336** (registered with PACCAR and serviced at a rooftop whose history was extracted), **PACCAR ONLY 724** (registered in the territory, no service history here) and **CDK ONLY 1,931** (serviced here, no PACCAR registration in this extract set). The last figure is the size of the all-makes and out-of-territory book and is the reason `v_rich_target` is smaller than the RO population. **The workbook is dealer payload and is never committed here** — the script is committed, the output is not.

```
python bin/build_outreach_workbook.py --db /path/to/cdk_wireframe.db \
    --out EVEglyphDesign_WarrantyGENE_Outreach_P1-P4_<date>.xlsx
```

Boundary

GENE fires **before** the repair order exists. PACCAR's PRWS remains authoritative on claim filing, and nothing in this model submits, adjudicates or authorises a claim. `v_rich_target` produces a prompt for a service writer to decide — a review filter, not a verdict.

Running it

```
python bin/build_wireframe_db.py \
  --src /dealer-owned/path/to/extracts \
  --db /dealer-owned/path/out/cdk_wireframe.db \
  --profile-out profile/$(date +%F)-extract-profile.json
```

The loader is idempotent, rebuilds from scratch each run, and prints an aggregate-only profile. It expects `P026.csv`, `P032.csv`, `P046.csv`, `P048.csv`, `Closed ROs.csv`, `Open ROs.csv`, `Customers.csv`, `service customer sales.csv`, `Customer total sales.csv`. Missing files are recorded as `SOURCE_MISSING` warnings rather than failing the run, because the extract set arrives in pieces. Powersports files sitting in the same directory are matched against `OUT_OF_SCOPE` and logged as `SCOPE_EXCLUDED`, never loaded.

SQLite is the wireframe target because it is inspectable with no server. The DDL is deliberately portable to the Azure PostgreSQL landing zone that the CDK Twin ingestion pipeline targets; see [DATA-QUALITY.md](#) for the promotion gates that must pass before a Postgres load is trustworthy.

Data Quality — 2026-09-17 extract set

Nothing is silently dropped. Every rejection lands in `dq_exception` with a rule code and a severity, and the counts below are reproducible by rerunning the loader.

RULE	SEVERITY	COUNT	MEANING
<code>VIN8_LENGTH</code>	WARN	1,223	RO rows whose <code>Vehicle ID</code> is not 8 or 17 characters, so they cannot reach the warranty registration
<code>DATE_PARSE</code>	WARN	213	PACCAR registration rows with a missing or unparseable in-service date — coverage end cannot be computed
<code>RO_DUPLICATE</code>	WARN	12	an RO number present in both the open and closed extract; the closed row wins
<code>SCOPE_EXCLUDED</code>	INFO	6	a file present in the extract directory that belongs to the BRP powersports lane and is deliberately not loaded — this wireframe is the truck business only
<code>CUST_UNKNOWN</code>	INFO	3	a customer on an RO absent from the customer master; a stub row is created

The five findings that constrain the model

1. Vehicle ID is a VIN8, not a VIN. 2,141 of 2,436 distinct CDK vehicle IDs are 8 characters, 245 are 6, and 4 are full 17-character VINs. The 6-character values cannot be lengthened by inference and must not be fuzzy-matched to a VIN8 — a wrong VIN match would attach one fleet's coverage to another fleet's truck. They are quarantined, counted, and excluded from `v_rich_target`, which is why the mart population is smaller than the RO population.

2. The four PACCAR reports are disjoint partitions, not overlapping views. P026 (587 VINs), P032 (273), P046 (196) and P048 (4) share zero VIN8 between them, and their city profiles differ — Fredericton and Dieppe in P026, Summerside and Cap Pelé in P032, Kentville and Truro in P046, Quebec towns in P048. They union to exactly 1,060 units, which matches the count of `BASE/VEH` coverage rows exactly: one base vehicle warranty per unit. Loading them as one union is therefore correct, and any attempt to reconcile them against each other is wasted work.

3. Coverage expiry has two axes and the mileage axis is already binding. Of 650 time-active coverages where an odometer reading exists, 170 are already past their mileage ceiling and 26 are within 25,000 of it. A time-only filter would present 170 dead coverages as live. `v_unit_coverage_state` exposes `is_time_active` and `miles_remaining_est` separately and never collapses them.

4. Owner name is not a join key. Only 65 of 534 manufacturer-side registered-owner names match a CDK RO customer name exactly. Registration owner is stored as an attribute of the unit and is used for outreach copy, never for identity resolution.

5. The aggregate extracts are not a population. `Customer total sales.csv` is a top-100 ranked report and `service customer sales.csv` covers 2,693 customers against 16,022 in the master. They are loaded into `agg_*` tables for tie-out and are barred from scoring, because a customer-period total cannot be attributed to a VIN.

Promotion gates — before any Azure PostgreSQL load is trusted

1. `dq_exception` carries zero `BLOCK` rows.
2. `SUM(total_sales) over ro_status='CLOSED'` ties to the closed-RO extract total (\$12,105,138.15 on this set, 7,534 ROs).
3. `COUNT(*)` in `fact_warranty_registration` equals the summed source report rows less `VIN8_LENGTH/BLOCK` rejections (3,258 on this set).
4. Every `fact_repair_order.vin8` that is non-null is 8 characters.
5. `v_rich_target` contains no row whose nearest expiry is negative.
6. `dim_unit` row count equals the distinct union of registered and serviced VIN8 (2,991 on this set).

Known gaps in the extract set

- No RO line detail — the extracts are RO headers only. Labour operation codes, SRT hours, part numbers and the 3-C complaint text are absent, so margin estimate and 3-C compliance cannot yet be scored per job. This is the single highest-value addition to the next extract request.
- No claim-side data on the truck lane. PRWS claim history is not in this set, so booked-versus-eligible cannot be closed on the Peterbilt side yet. This is the second highest-value addition to the next extract request. The 205 Warranty on Demand claim records in the extract directory do not fill this gap — they are BRP powersports claims and are excluded by scope. The PACCAR Warranty Procedure Manual v2026.7 is present and is policy rather than data; it governs how a claim is filed, not what is in the file.
- No SmartLINQ / PACCAR Solutions fault-code feed, so the event-driven rung is absent.
- Odometer arrives only via ROs, so units never serviced here have no mileage reading and their mileage axis is unknown — reported as `NULL`, never as zero.